

Figure 3. Two kernels held 77 percent of decode time.

Measured over 28.85 ms of GPU work per decoded token on an L40S. CUDA graph capture had already removed launch overhead, so the 2,236 elementwise launches cost only 3.4 ms between them and kernel count was not the bottleneck. This measurement is what redirected the work from reducing launches to reshaping two kernels. Source: engine/klinear/DECODE-PROFILE.md

