

Figure 5. What INT4 costs, measured against the BF16 checkpoint.

Both sides served by the same vLLM build on one H200 with identical prompts, so only the weights differ. Keeping the shared experts in BF16 improves every agreement metric shown yet makes perplexity worse by 1.95 percent, which is why it was not adopted. These quantisation costs are an order of magnitude larger than the kernel differences quoted in Figure 1. Source: `engine/accuracy/RESULTS.md`

